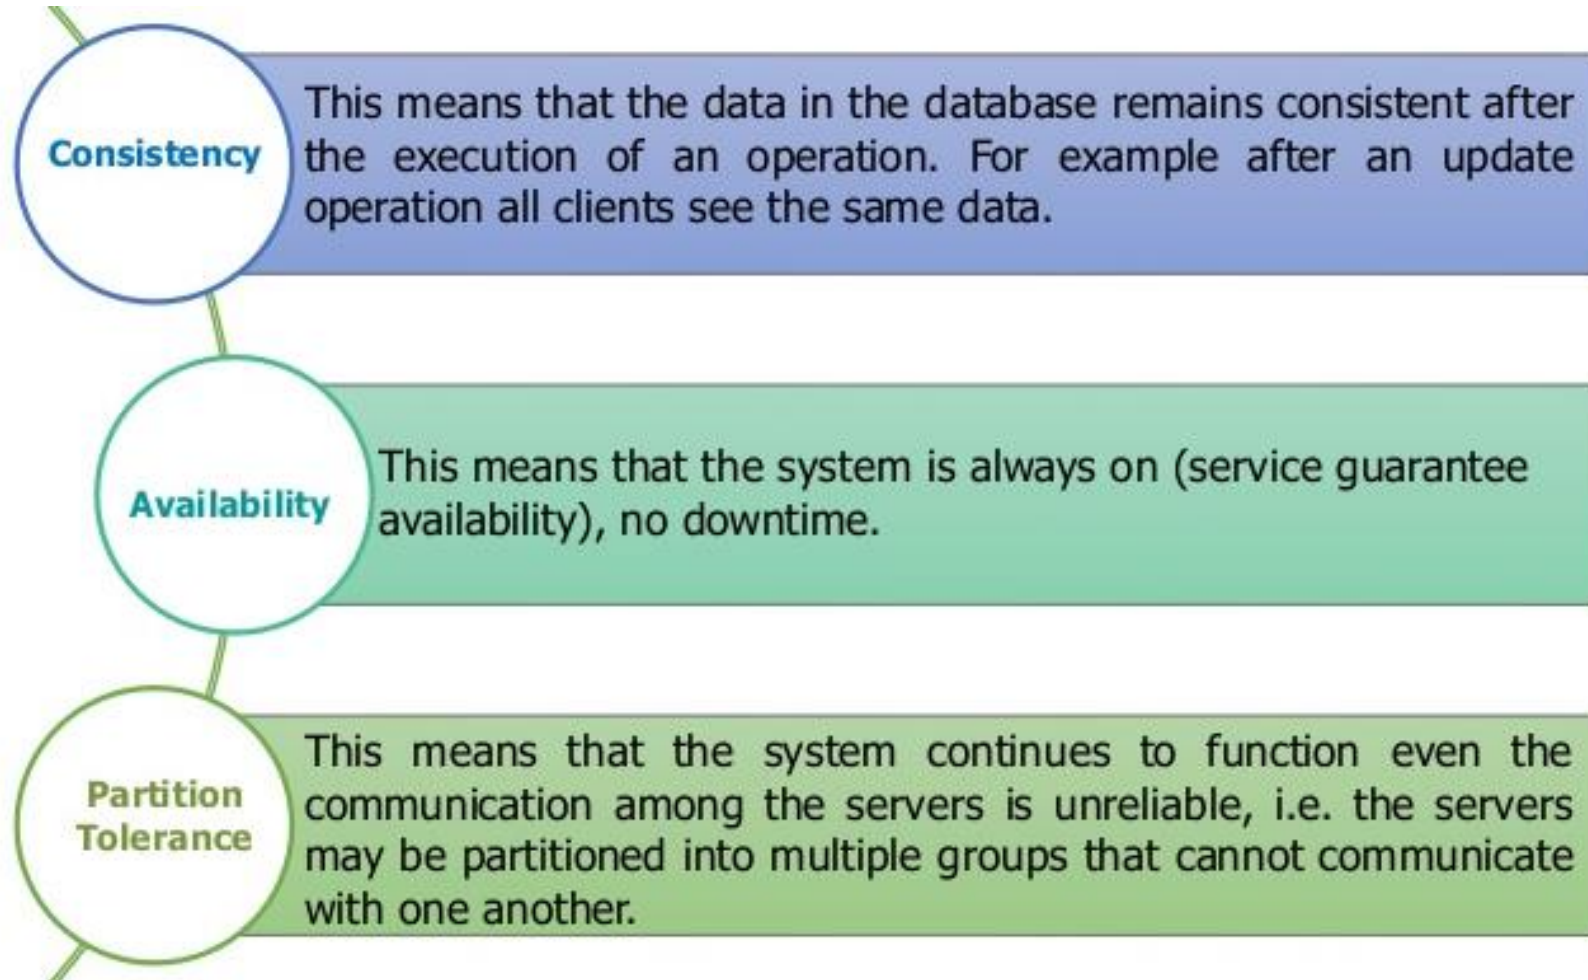


# NoSQL Databases: HBase

# Data Architecture

- No standard solution that fits all
- Business and data defines the architecture
- Multiple databases, different types depending on the characteristics of each data subset

# CAP



# CAP

- **Consistency** - A read is guaranteed to return the most recent write for a given client.
- **Availability** - A non-failing node will return a reasonable response within a reasonable amount of time (no error or timeout).
- **Partition Tolerance** - The system will continue to function when network partitions occur.


# CAP Theorem

**CA** - Single site cluster, therefore all nodes are always in contact. When a partition occurs, the system blocks.


**CP** - Some data may not be accessible, but the rest is still consistent/accurate.

**AP** - System is still available under partitioning, but some of the data returned may be inaccurate.

CP



AP



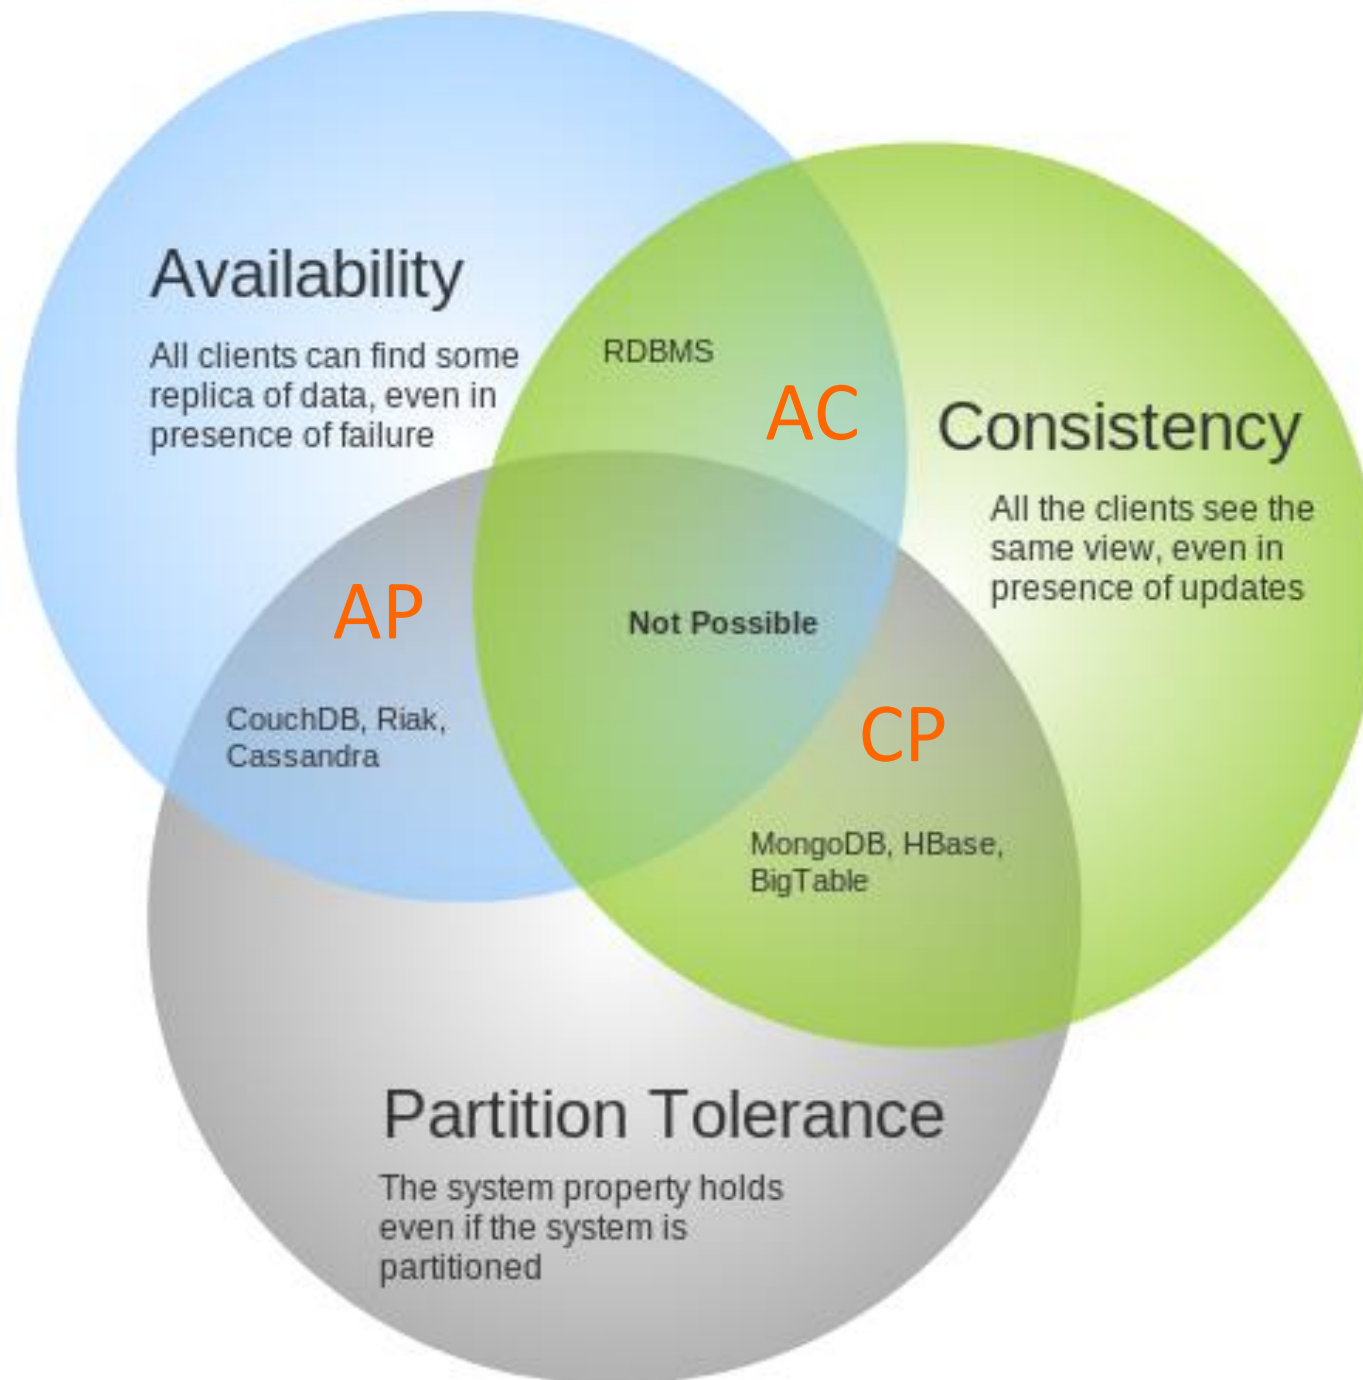
# CAP Theorem

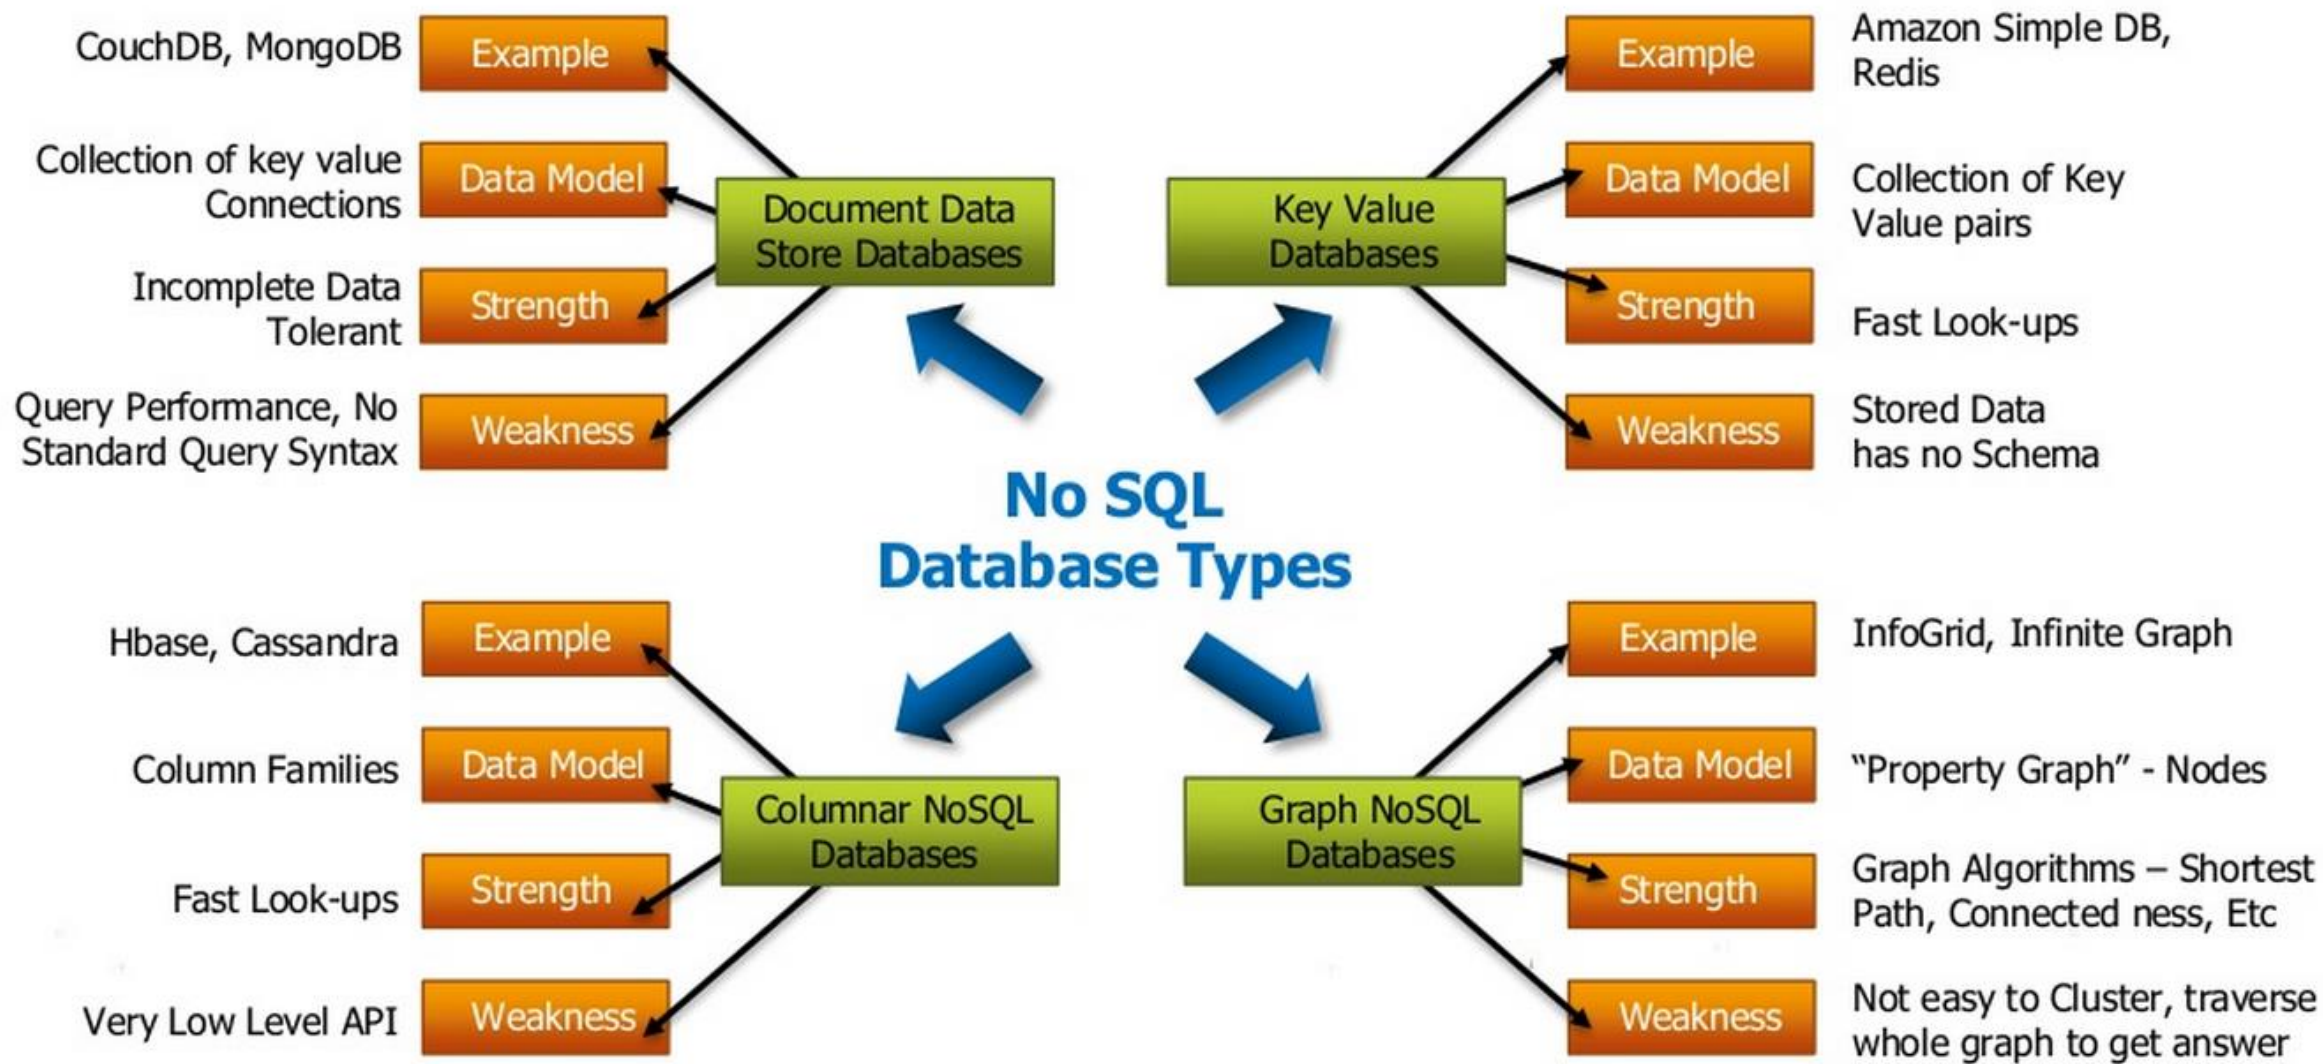
CAP provides the basic requirements for a distributed system to follow **2 of the 3 requirements**.

In theoretically it is **impossible** to fulfill all 3 requirements.

Therefore all the current NoSQL database follow the different **combinations of the C, A, P** from the CAP theorem.







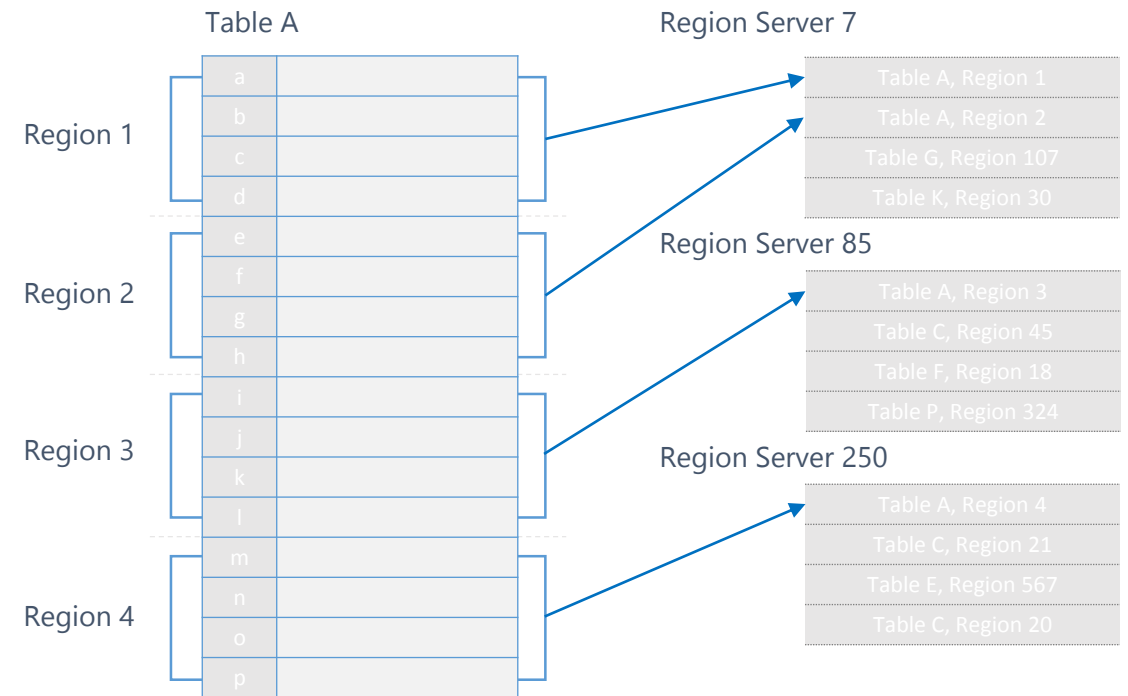
# What is HBase

- Distributed, non-relational database
  - Columnar, schema-free data model
  - NoSQL on top of Hadoop
- Large scale
  - Linear scalability
  - Billions of rows X millions of columns
  - Many deployments with 1000+ nodes, PBs of data
- Low latency
  - Real-time random read/writes
- Open Source
  - Modeled after Google's BigTable
  - Started in 2006



# Data Model

- Scale-out architecture
  - Automatic sharding of tables
  - Automatic failover
  - Strong consistency for reads and writes
- APIs
  - Get/Put
  - Scan
  - Coprocessors



# Performance Features

- Column Families
- In-memory caching
- High throughput streaming writes

Row Key	Customer		Sales	
Customer Id	Name	City	Product	Amount
101	John White	Los Angeles, CA	Chairs	\$400.00
102	Jane Brown	Atlanta, GA	Lamps	\$200.00
103	Bill Green	Pittsburg, PA	Desk	\$500.00
104	Jack Black	St. Louis, MO	Bed	\$1600.00

Column Families

# Sharding

- Holding rows of database on different partitions
- Same table divided onto different servers, even different geographies
- Reduces index size

# Sharding

- More reliance on interconnection between servers
- Increased latency in querying when more than one shard must be searched
  - Some searches are fast, others are slow
- Often no guarantees about cross shard consistency

# Notable Capabilities

- Integration features
  - Integration with Hadoop MapReduce, Hive, Tez
  - Bulk import of large amounts of data
- Client APIs
  - Java, REST, python, node.js, php, .NET

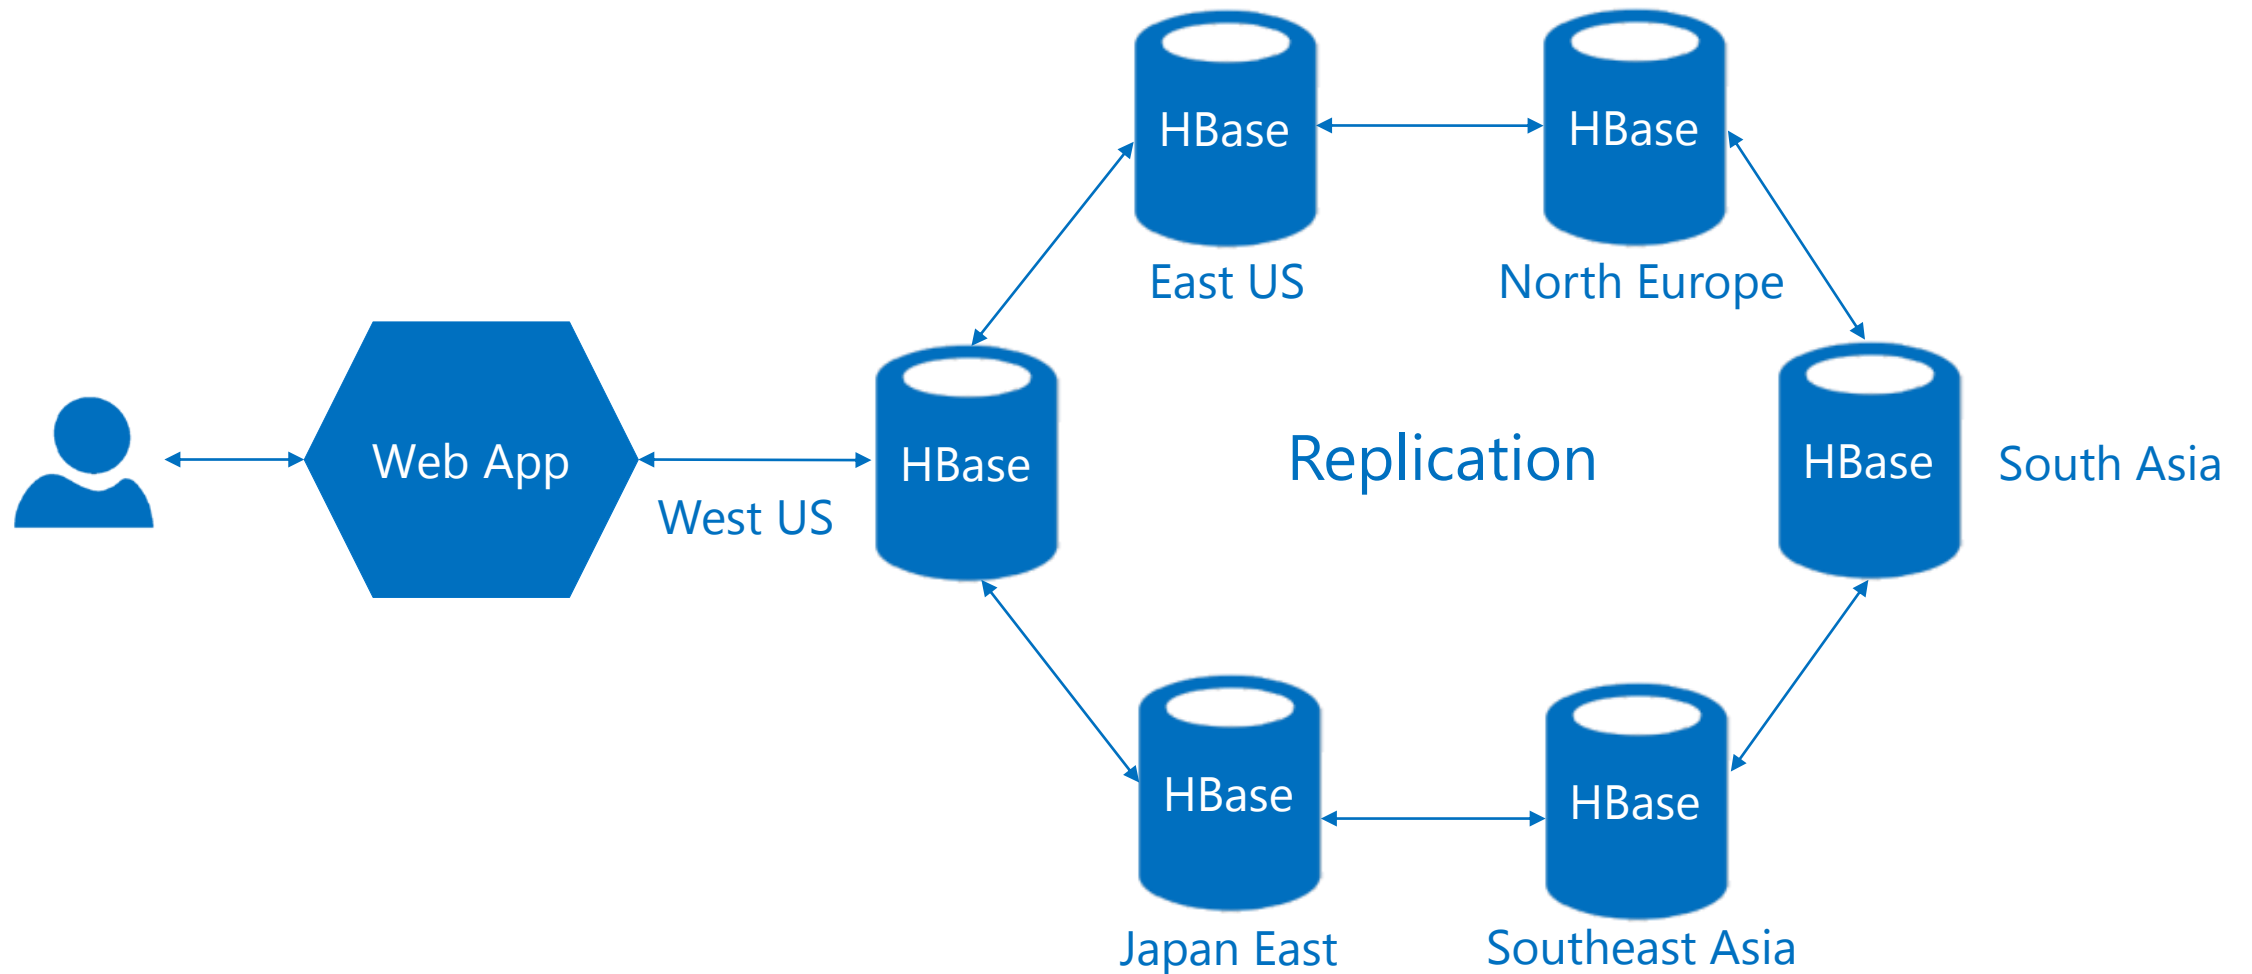


# Use case #1: key value store

- Key value store
  - Message systems
  - Content management systems
- Examples
  - Facebook Messages
  - Twitter-like messages
  - Webtable – web crawler/indexer




# Use case #1: key value store




# Use case #2: sensor data

- Sensor data
  - Social analytics
  - Time series databases
  - Interactive dashboards with trends, counters, etc
  - Audit log systems



# Use case #2: sensor data



# HBase as a platform

- Running on top of HBase using it as a datastore:
  - Phoenix                      OpenTSDB
  - Kiji                              Tephra
  - Titan                             Kylin
- Integrated with HBase:
  - Hive                              Pig
  - Storm                             Spark
  - Flume                            Solr
  - Ganglia