

Setting Up a Relational Database in the Cloud

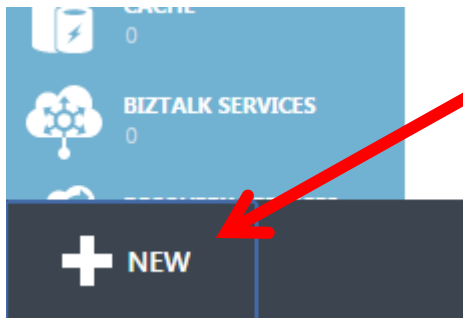
Contents

Exercise 1: Provisioning an Azure SQL Database.....	3
Exercise 2: Configuring DB Firewall and Access Rights	6
Exercise 3: Loading Data Tables into the Database	9
Exercise 4: Reading from Azure SQL Database.....	12
Exercise 5: Forming a Relational Database.....	13

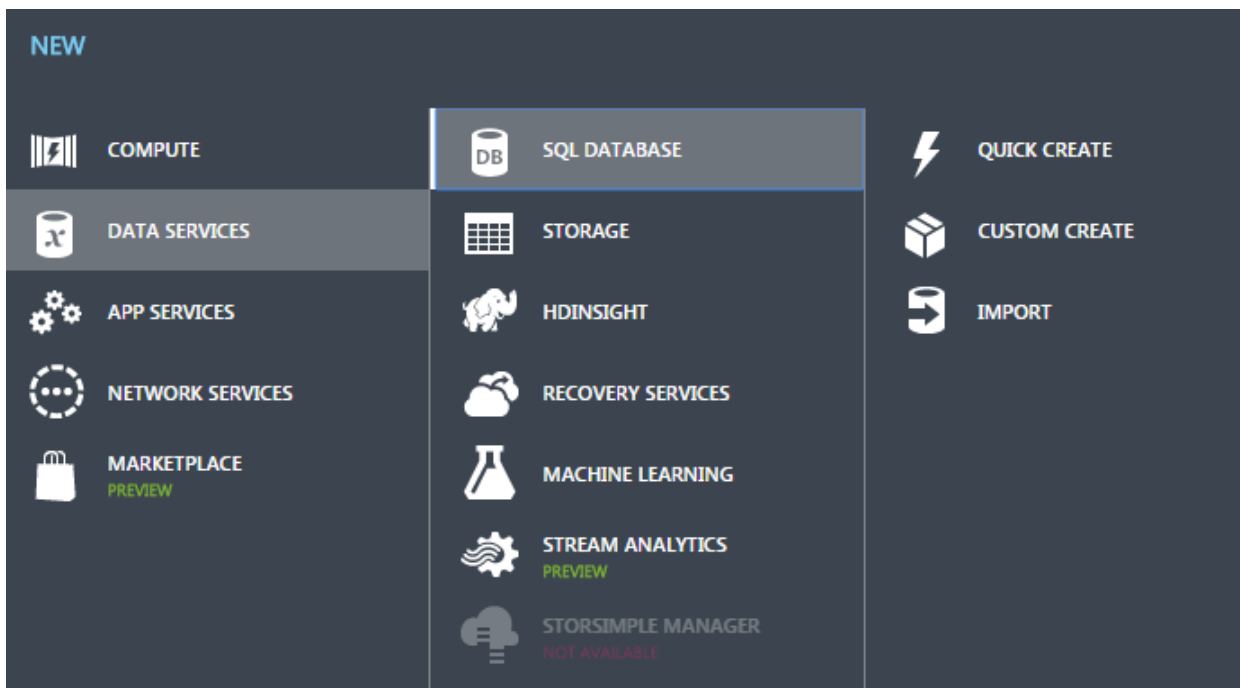
Exercise 1: Provisioning an Azure SQL Database

In the SaaS era of cloud computing, databases are now commoditized. Cloud-based databases are managed free and elastically scaled for maximum efficiency. This frees up DBA staff and allows companies to refocus on the core of the business.

1. Visit the Azure Management portal:
 - a. <https://manage.windowsazure.com>
2. Provisioning an Azure SQL DB
 - a. Within the Azure Management portal, select **new**.



- b. **New > Data Services > SQL Database > Custom Create**



3. Set Database Specifications

NEW SQL DATABASE - CUSTOM CREATE

Specify database settings

NAME

SUBSCRIPTION
Free Trial (5576a6d9-c6fe-45a3-a834-9fc58b7c) ▼

SERVICE TIERS
BASIC STANDARD PREMIUM

RETIRED TIERS
WEB BUSINESS ?

PERFORMANCE LEVEL
S0 (10 DTUs) ▼ ?

MAX SIZE
250 GB ?

COLLATION
SQL_Latin1_General_CP1_CI_AS ▼ ?

SERVER
Choose a server ▼

- Name:** Name the database. A database is a collection of tables. A server may have multiple databases on it.
- Subscription:** The subscription that the database will be charged to.
- Server Tiers:** The quality of the server. Choose **basic tier** for this example. Visit the following link for a comparison chart of the 3 tiers. Do not choose web or business tiers; they will be phased out in the near future. <http://azure.microsoft.com/en-us/pricing/details/sql-database/>
- Performance Level:** Select "**Basic (5 DTUs)**" for this lab. Specify the workload level of the database. Visit the link in server tiers for a full comparison chart of each performance level.
- Collation:** Leave this on default (**SQL_Latin1_General_CP1_CI_AS**). This is the character encoding within the table data fields. For additional information, visit: <https://msdn.microsoft.com/en-us/library/azure/ee336245.aspx#sscs>
- Server:** Choose **new server**. This server will host all your databases. A server name will be automatically assigned.

4. Server Credentials

CREATE SERVER

SQL database server settings

LOGIN NAME

LOGIN PASSWORD

CONFIRM PASSWORD

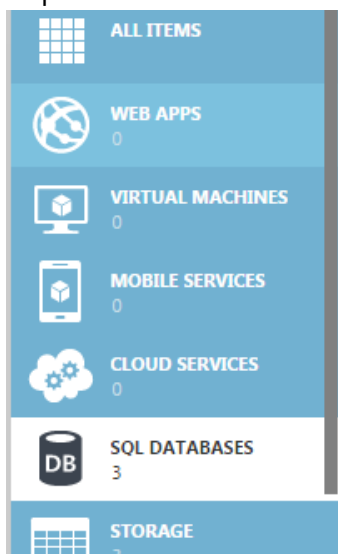
REGION

South Central US ▼

☒ ALLOW WINDOWS AZURE SERVICES TO ACCESS THE SERVER. ?

☒ ENABLE LATEST SQL DATABASE UPDATE (V12) ?

- Login Name:** The login name of the server. All the databases in the server will require this login credential.
 - Password:** Password of the server.
 - Region:** For the purposes of this lab, choose **West Europe**. This is the region in which the database will be hosted. Be sure to keep the database server in the same region as your servers, clients, or other data pipeline artifacts.
 - Allow Windows Azure Services to Access the Server:** Check this box. This will allow other Azure Services such as Stream Analytics to link with the database easily.
 - Enable Latest SQL Database Update (v12):** Checked.
5. Confirm the existence of the server in the Azure Management Portal. The database will take ~2 minutes to provision.



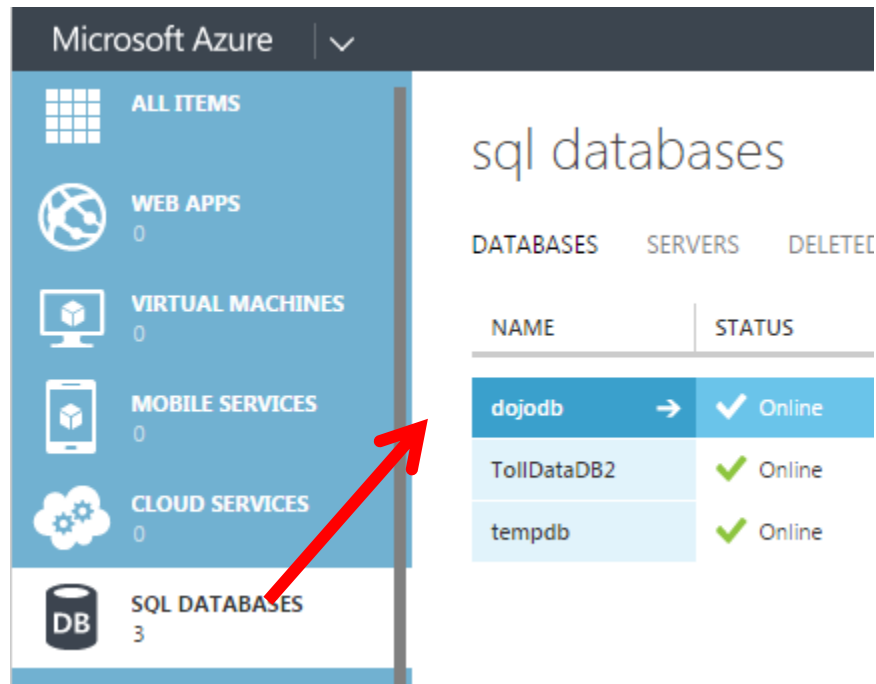
sql databases

DATABASES SERVERS DELETED DATABASES

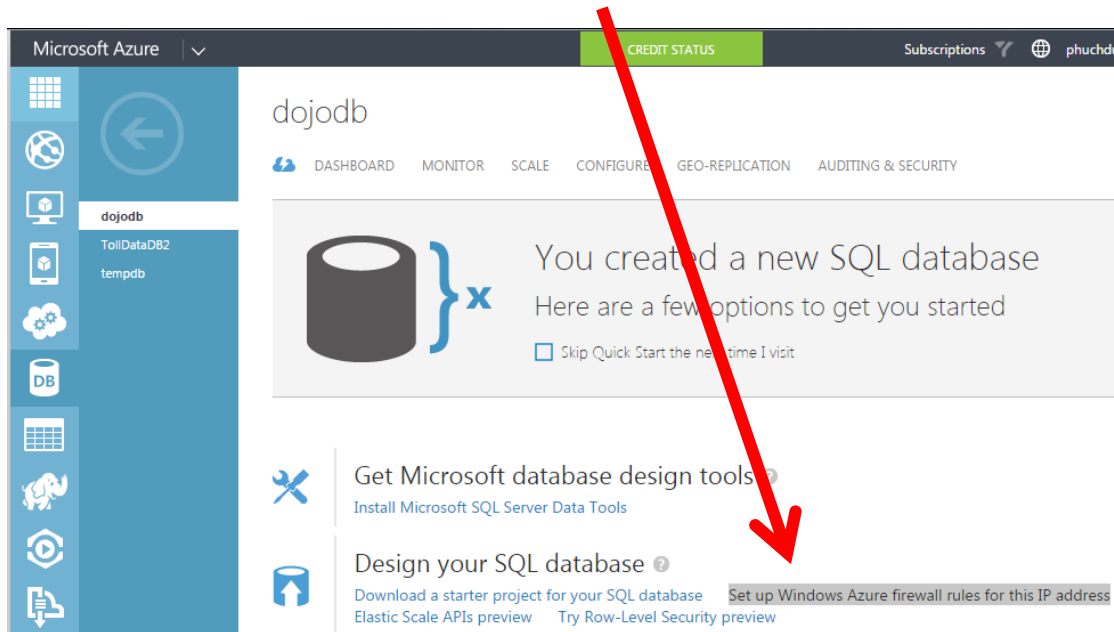
NAME	STATUS	REPLICATION
dojodb	✓ Online	None
TollDataDB2	✓ Online	None
tempdb	• • Creating	None

Exercise 2: Configuring DB Firewall and Access Rights

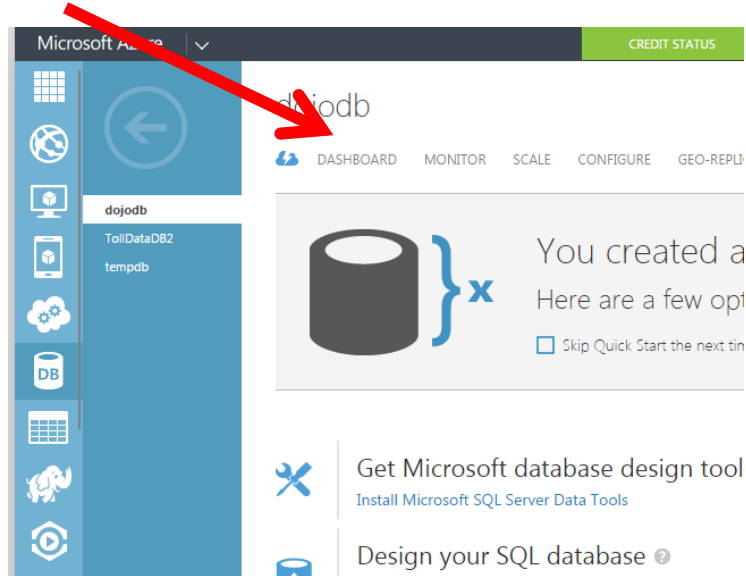
1. Add your current computer IP address to the firewall rules.
 - a. Azure Management Portal > SQL Databases > YourDB



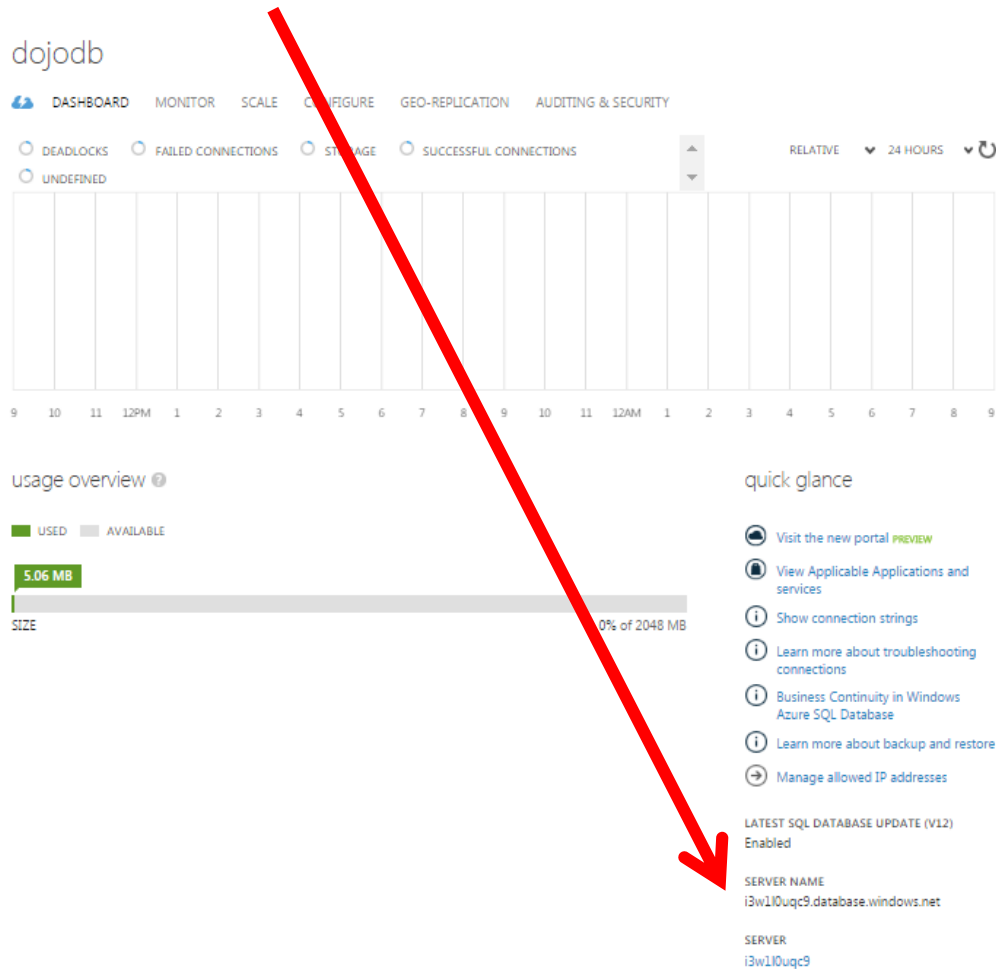
- b. Select "Set up Windows Azure firewall rules for this IP address" to automatically add the IP address of the current device to the database's firewall rules.



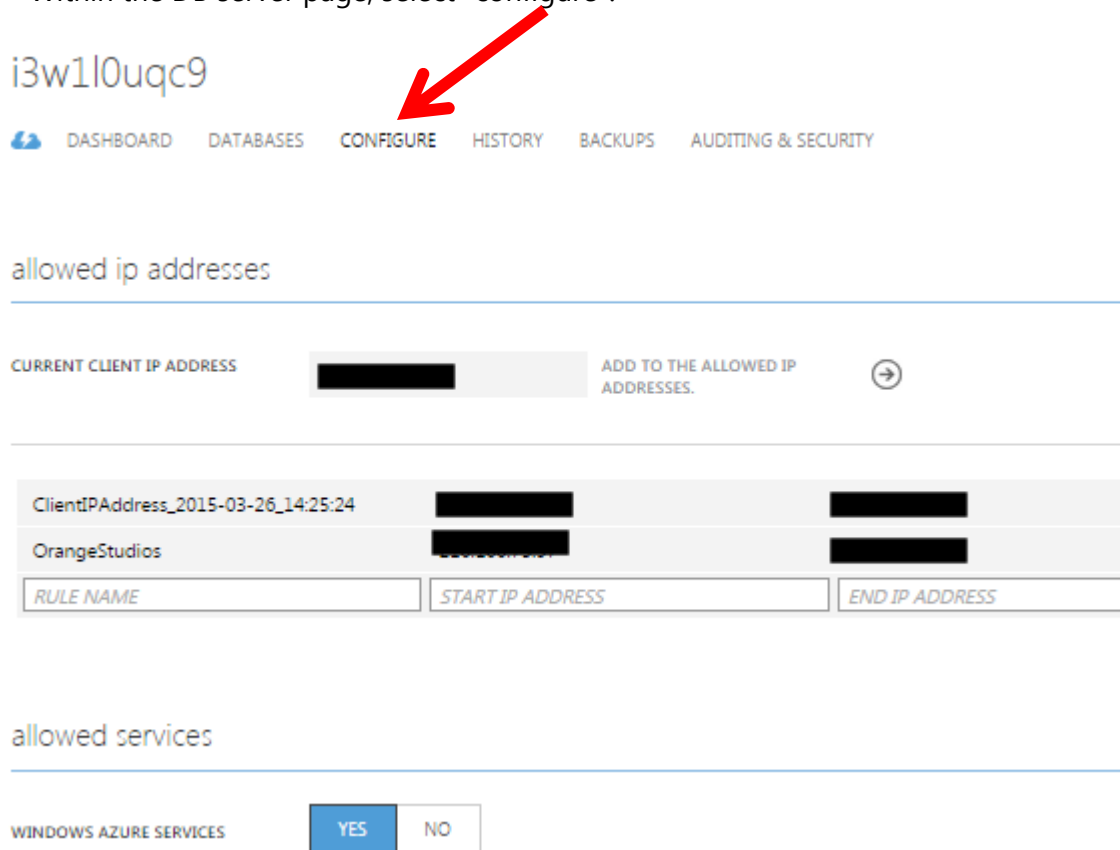
2. View firewall rules:
 - a. Select "Dashboard" from the DB view.



- b. Click on the DB's server name within the dashboard.



- c. Within the DB server page, select "configure".



i3w1l0uqc9

[DASHBOARD](#) [DATABASES](#) [CONFIGURE](#) [HISTORY](#) [BACKUPS](#) [AUDITING & SECURITY](#)

allowed ip addresses

CURRENT CLIENT IP ADDRESS ADD TO THE ALLOWED IP ADDRESSES. [➔](#)

ClientIPAddress_2015-03-26_14:25:24	<input type="text" value=""/>	<input type="text" value=""/>
OrangeStudios	<input type="text" value=""/>	<input type="text" value=""/>
RULE NAME	START IP ADDRESS	END IP ADDRESS

allowed services

WINDOWS AZURE SERVICES ☒ YES ☐ NO

- d. This screen allows users to manually add IP addresses, in case the company or router has an external IP as well.
- e. Be sure that the "Windows Azure Services" is enabled with "YES".

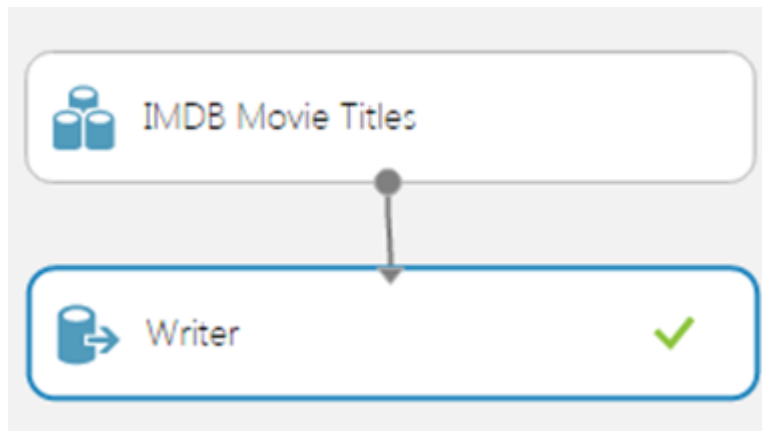
Exercise 3: Loading Data Tables into the Database

There are many methods to connect to the database. Some include:

- Azure PowerShell
- Azure Machine Learning Studio
- MS Access
- MS Excel
- Programmatically (sqlcmd, ASP.net, ODBC, PHP, JDBC, ADO.NET)
- SQL Server Management Studio (SSMS)
- Visual Studio

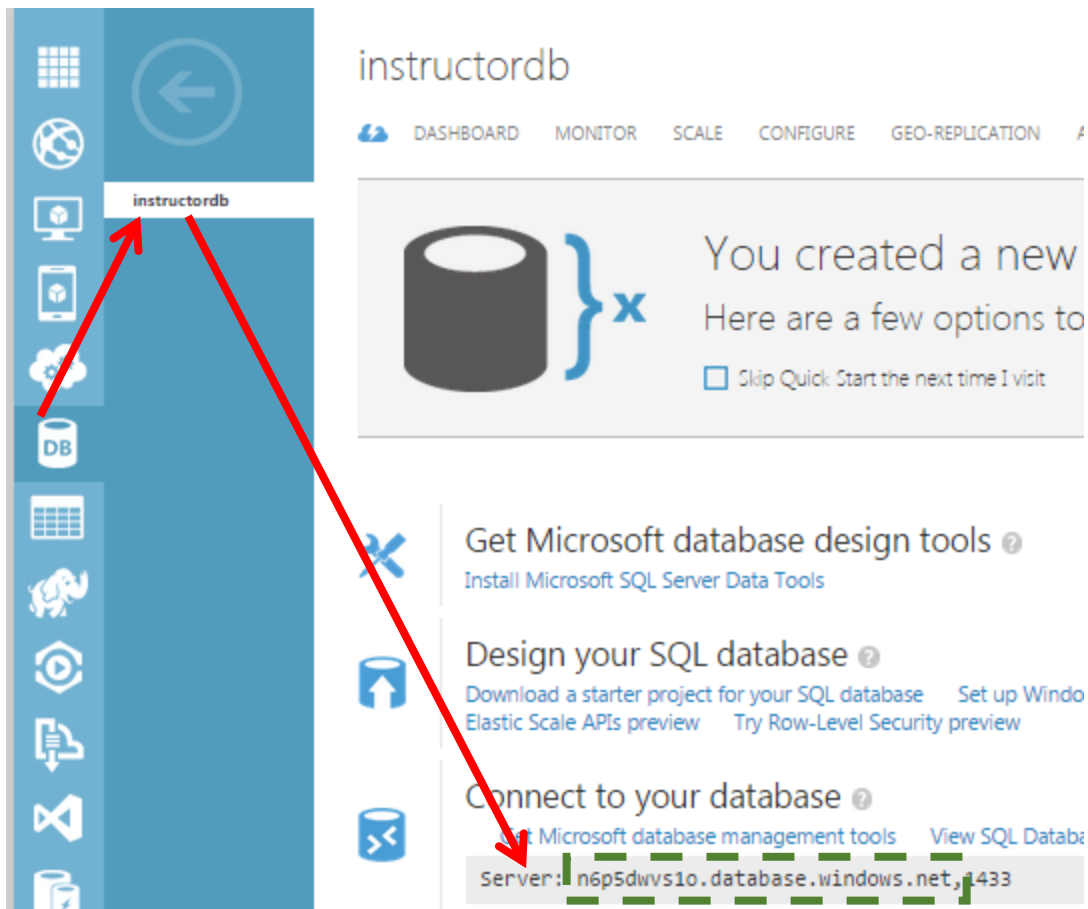
This exercise will focus on connecting to an Azure SQL DB using Azure Machine Learning Studio.

1. Log into Azure ML and create a new experiment.
 - a. <https://studio.azureml.net>
2. The datasets.
 - a. A database is just a collection of tables which may reference each other in a “key” and index fashion.
 - b. Drag in the following three datasets: **IMDB Movie Titles, Movie Ratings, Movie Tweets**
3. Writing to Azure SQL DB: The Azure ML writer module
 - a. Drag in a writer module and connect it to “IMDB Movie Titles”
 - b. The Writer Module will dynamically create schemas and specify the correct datatypes for each of the columns. No coding or SQL necessary.



c. Writer Credentials:

- i. **Please specify data destination:** Azure SQL Database
- ii. **Database server name:** <YourServerName>.database.windows.net
 - (1) Within your Azure Management Portal, select SQL Databases > YourDB
 - (2) Find the "Server" address under "Connect to your Database"



- iii. **Database name:** the name of your database. In the example above, the database is called instructordb.
- iv. **Server user account name:** the server's user name that was set by the user during provisioning.
- v. **Server user account password:** the server's login password that was set by the user during provisioning.
- vi. **Accept any server certificate (insecure):** unchecked
- vii. **Comma separated list of columns to be saved:** these are the columns to be saved. List the column names in order, separated by commas. In the example above, the dataset "IMDB Movie Titles" has the following columns to be saved: 'Movie ID, Movie Name'
- viii. **Data table name:** this is what the table will be called. Call this table 'movies'.
- ix. **Comma separated list of data table columns:** this is how the columns will be provisioned within the database table. This is a chance to rename the columns. In this example, fill in "filmid, moviename". Spaces in column names will break the database. Also, it's good style not to include capitalizations within column names as it will improve the efficiency when writing queries.
- x. **Number of rows written per SQL Azure operation:** 50

4. Run the experiment to write to the database. This will take ~4 minutes to write 16,000 rows to the database.

The screenshot displays the Data Science Dojo interface. On the left, a workflow is shown with two components: 'IMDB Movie Titles' (represented by a database icon) and 'Writer' (represented by a database icon with a green checkmark). An arrow points from 'IMDB Movie Titles' to 'Writer'. On the right, the 'Writer' configuration panel is open, showing various settings for writing data to a database.

Writer

Please specify data destination
Azure SQL Database

Database server name
n6p5dwvs1o.database.windows

Database name
instructordb

Server user account name
dojoadmin

Server user account password

☐ Accept any server certificate

Comma separated list of columns
Movie ID, Movie Name

Data table name
movies

Comma separated list of delimiters
filmid, filename

Number of rows written per batch
50

Exercise 4: Reading from Azure SQL Database

The reader module allows users to query the database.

1. Create a new experiment.
2. Drag in a reader module.
3. Reader module credentials:
 - a. **Data source:** Azure SQL Database
 - b. **Database server name:** <YourServerName>.database.windows.net
 - c. **Database query:** the following query returns the first 10 rows of the table.

```
select top 10 *  
from movies;
```

Final result:

The screenshot shows the Data Science Dojo interface. On the left, a canvas displays a 'Reader' module with a green checkmark, indicating it is successfully configured. On the right, the 'Reader' module configuration panel is visible. It includes fields for 'Data source' (Azure SQL Database), 'Database server name' (n6p5dwvs1o.database.windows.net), 'Database name' (instructordb), 'Server user account name' (dojoadmin), and 'Server user account password' (masked with asterisks). There is also a checkbox for 'Accept any server certificate (insecure)'. The 'Database query' field contains the SQL query: `1 select top 10 *
2 from movies;`

4. Run the experiment, and visualize the output.


rows: 10, columns: 3

	id	filmid	filmname
view as			
	231	8	Edison Kinetoscopic Record of a Sneeze (1894)
	232	91	Le manoir du diable (1896)
	233	417	Le voyage dans la lune (1902)
	234	628	The s of Dollie (1908)
	235	833	The Country Doctor (1909)
	236	1223	Frankenstein (1910)
	237	1740	The Lonedale Operator (1911)
	238	2101	Cleopatra (1912)
	239	2130	Linferno (1911)
	240	2844	Fantmas - l'ombre de la guillotine (1913)

Exercise 5: Forming a Relational Database

Right now, the database has a single table, which is just a database table. For this database to truly be a relational database, other tables that reference each other in relationships must also be present. In this exercise, we will write additional tables to the database, "movie ratings" and "movie tweets", which will take 7 and 10 minutes to write to the database respectively. To increase the write times, consider sampling the datasets first using the sampling module.





1. Write the dataset called '**Movie Ratings**' to the database. 227472 rows by 4 columns.
 - a. Comma separated list of columns to be saved: '**UserId, MovieId, Rating, Timestamp**'
 - b. Data Table Name: '**ratings**'
 - c. Comma separated list of data table columns: '**userid, filmid, rating, timestamp**'
 - d. Write time: ~7 minutes
2. Write the dataset called '**Movie Tweets**' to the database. 170285 rows by 8 columns.
 - a. Comma separated list of columns to be saved: '**Scraping Time, Tweet ID, User ID, Movie ID, Rating, Retweet Count, Favorite Count, Time Zone**'
 - b. Data Table Name: '**tweets**'
 - c. Comma separated list of data table columns: '**scrapetime, tweetid, userid, filmid, rating, retweetcount, favoritecount, timezone**'
 - d. Write time: ~10 minutes



3. View the database
 - a. Create a new experiment.
 - b. Drag in a reader module.
 - i. The following query will list all the tables inside the database. Be sure that the table 'movies', 'ratings', and 'tweets' are in the database.

```
SELECT * FROM
information_schema.tables
```










ii. Run the experiment and visualize the output.

TABLE_CATALOG	TABLE_SCHEMA	TABLE_NAME	TABLE_TYPE
			
instructordb	sys	database_firewall_rules	VIEW
instructordb	dbo	iris	BASE TABLE
instructordb	dbo	city	BASE TABLE
instructordb	dbo	movies	BASE TABLE
instructordb	dbo	ratings	BASE TABLE
instructordb	dbo	tweets	BASE TABLE

c. Drag in a second reader module.

i. Probe the tweets table.

SELECT TOP 10 * FROM tweets;

id	scrapetime	tweetid	userid	filmid	rating	retweetcount	favoritecount	timezone
								
1	1391030896	307139025897152000	296041028	444778	8	0	0	Bangkok
2	1391030896	307139922186366000	239615275	1411238	6	0	0	Eastern Time (US & Canada)
3	1391030895	307142713369849000	36937267	1496422	7	0	0	Belgrade
4	1391030895	307143340607041000	45826125	118799	5	0	0	Budapest
5	1391030894	307144290214215000	45826125	338013	4	0	0	Budapest
6	1391030894	307144402382491000	27662592	1003052	8	0	0	Prague
7	1391030893	307144992542031000	45826125	64665	8	0	0	Budapest
8	1391030893	307145230707220000	206144706	884328	7	0	0	Riyadh
9	1391030892	307145590041624000	114473716	1559547	9	0	0	Taipei
10	1391030891	307149804595716000	124670134	949731	5	0	0	Unknown

- d. Drag in a third reader module.
- i. Probe the ratings table.

```
SELECT TOP 10 * FROM ratings;
```

id	userid	filmid	rating	timestamp
1	1	68646	10	1381620027
2	1	113277	10	1379466669
3	2	454876	8	1394818630
4	2	790636	7	1389963947
5	2	816711	8	1379963769
6	2	1091191	7	1391173869
7	2	1322269	7	1391529691
8	2	1433811	8	1380453043
9	2	1454468	8	1387016442
10	2	1535109	8	1386350135

Final output:

Reader
List of Tables
 ✓ ^

Reader
Movies
 ✓ ^

Reader
Ratings
 ✓ ^

Reader
Tweets
 ✓ ^






4. Perform a join on the 3 tables
 - a. Drag in another reader module and try to join the tables.

b. Below is a sample code on joining two tables together.

```
select
    m.filmname as Film,
    t.userid as UserID,
    t.rating as UserRating,
    t.timezone as Region
from tweets as t
join movies as m
    on t.filmid = m.filmid
```

rows
338404

columns
4

	Film	UserID	UserRating	Region
view as 				
	Mah nakorn (2004)	296041028	8	Bangkok
	Mah nakorn (2004)	296041028	8	Bangkok
	No Strings Attached (2011)	239615275	6	Eastern Time (US & Canada)
	No Strings Attached (2011)	239615275	6	Eastern Time (US & Canada)
	The Paperboy (2012)	36937267	7	Belgrade
	The Paperboy (2012)	36937267	7	Belgrade
	La vita ? bella (1997)	45826125	5	Budapest
	La vita ? bella (1997)	45826125	5	Budapest
	Eternal Sunshine of the Spotless Mind (2004)	45826125	4	Budapest
	Eternal Sunshine of the Spotless Mind (2004)	45826125	4	Budapest